

# Self-Driving Car Engineer Nanodegree

## Deep Learning

### Project: Build a Traffic Sign Recognition Classifier

For this project, we can see each of the following setups as outlined below, for necessary completion of the rubric points. The next 5 sections expound on these details.

1. Load the data set
2. Explore, summarize and visualize the data set
3. Design, train and test a model architecture
4. Use the model to make predictions on new images
5. Analyze the softmax probabilities of the new images
6. Summarize the results with a written report

This here Jupyter notebook is the written report!

---

## Step 0: All needed imports

Useful for testing if environment is ready to rip

```
In [1]: import csv
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import numpy as np
import pickle
import random

from sklearn.utils import shuffle
from sklearn.model_selection import train_test_split

import tensorflow as tf
from tensorflow.contrib.layers import flatten

print('Loaded all imports!')

Loaded all imports!
```

---

## Step 1: Load the data set

Retrieved file via curl piping into a zip file and extracting. Per project description, [dataset may be found at the provided link](https://d17h27t6h515a5.cloudfront.net/topher/2017/February/5898cd6f_traffic-signs-data/traffic-signs-data.zip) ([https://d17h27t6h515a5.cloudfront.net/topher/2017/February/5898cd6f\\_traffic-signs-data/traffic-signs-data.zip](https://d17h27t6h515a5.cloudfront.net/topher/2017/February/5898cd6f_traffic-signs-data/traffic-signs-data.zip)).

```
In [2]: # Load pickled data

training_file = './train.p'
validation_file = './valid.p'
testing_file = './test.p'

# sign names for labels
names_file = './signnames.csv'

with open(names_file) as f:
    reader = csv.reader(f)
    next(reader) # skip first row
    names = {int(rows[0]):rows[1] for rows in reader}

with open(training_file, mode='rb') as f:
    train = pickle.load(f)
with open(validation_file, mode='rb') as f:
    valid = pickle.load(f)
with open(testing_file, mode='rb') as f:
    test = pickle.load(f)

X_train, y_train = train['features'], train['labels']
X_valid, y_valid = valid['features'], valid['labels']
X_test, y_test = test['features'], test['labels']

print('all done loading')

all done loading
```

---

## Step 2: Dataset Summary & Exploration

The next few section provide a snapshot of the data and visualizes some examples.

### Here is Basic Summary of the Data Set

```
In [3]: # Number of training examples
n_train = np.shape(X_train)[0]

# Number of validation examples
n_validation = np.shape(X_valid)[0]

# Number of testing examples.
n_test = np.shape(X_test)[0]

# What's the shape of an traffic sign image?
image_shape = np.shape(X_train[0])

# How many unique classes/labels there are in the dataset.
sign_classes, class_counts = np.unique(y_train, return_counts = True)
num_classes = len(set(y_train)) # Original way I came up with, found t
his way above when needing to do augmentation

print("Number of training examples =", n_train)
print("Number of testing examples =", n_test)
print("Image data shape =", image_shape)
print("Number of classes =", len(sign_classes))

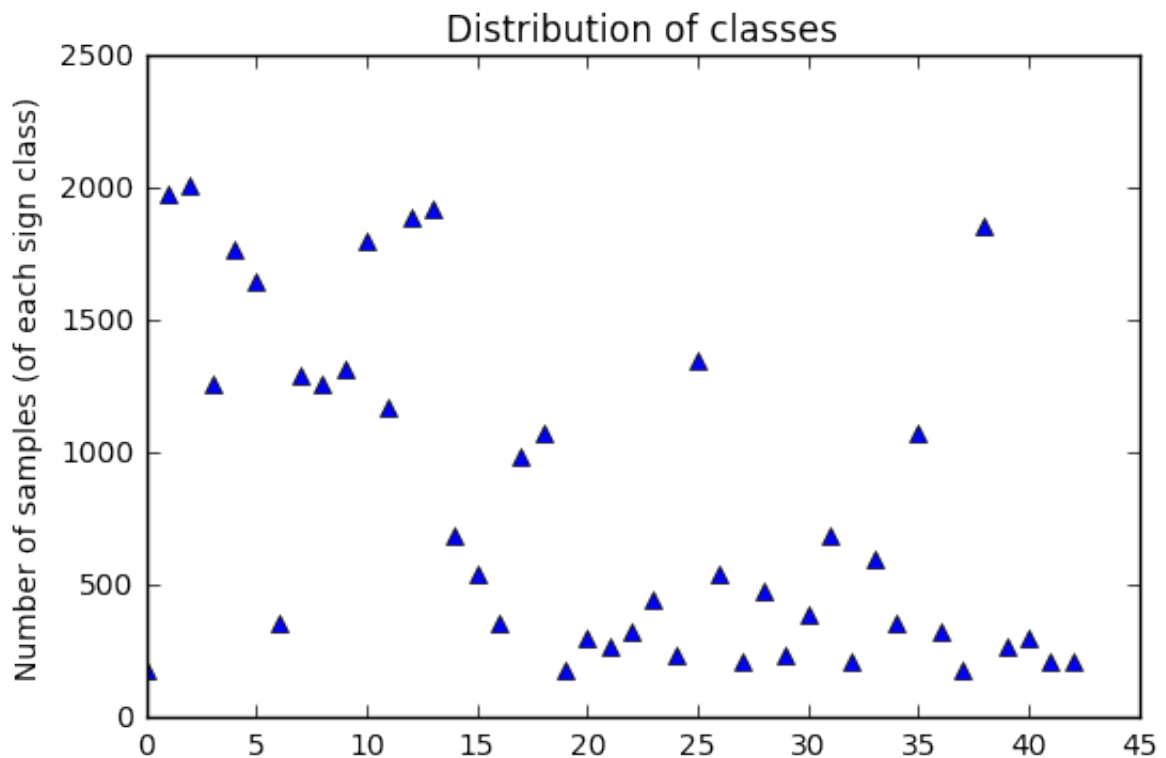
Number of training examples = 34799
Number of testing examples = 12630
Image data shape = (32, 32, 3)
Number of classes = 43
```

## here is an exploratory visualization of the dataset

As can be clearly seen, the distribution of classes is very uneven

```
In [88]: %matplotlib inline

plt.title("Distribution of classes")
plt.ylabel('Number of samples (of each sign class)')
chart = plt.plot(class_counts, 'b^')
#chart.set_xticklabels(names) #TODO Actually name the signs on the bar
chart so we know what is underrepresented\
#plt.bar(x, y, width, color="blue")
plt.show()
```



### ...and finally, a sampling of some images from the dataset

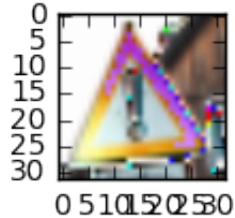
I ran this after shuffling and normalizing the dataset. Hence no need to randomize index. Commented in case needed

```
In [114]: %matplotlib inline

for x in range(5):
    index = x #random.randint(0, len(X_train))
    image = X_train[index].squeeze()

    plt.figure(figsize=(1,1))
    plt.title(names[y_train[index]])
    #plt.subplot(1,10,x+1) # failed attempt to inline some of the images, not germane to this work
    plt.imshow(image) # showing with cmap='gray' does not have effect; per docs, if array depth is 3, then cmap is ignored
    plt.show()
    print(y_train[x])
```

General caution



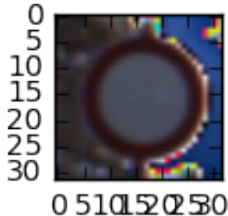
18

Keep right



38

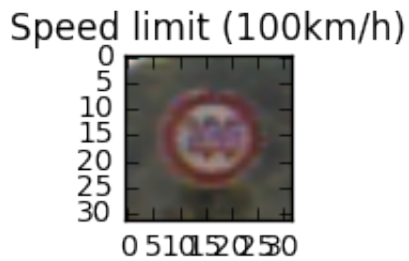
No vehicles



15



22



7

## Step 3: Design, train and test a model architecture

Here we setup the model architecture, pre-process the data and other fun things.

### Pre-process the Data Set (normalization, grayscale, etc.)

```
In [4]: X_train, y_train = shuffle(X_train, y_train)
X_train, X_validation, y_train, y_validation = train_test_split(X_train,
y_train, test_size=0.2, random_state=0)

print('shuffled up and training/validation sets ready')

shuffled up and training/validation sets ready
```

```
In [5]: def normalize(pixels):
        return (pixels - 128) / 128

X_train = normalize(X_train)
X_validation = normalize(X_validation)
X_test = normalize(X_test)

print('Images are now roughly normalized, hooray!')
```

Images are now roughly normalized, hooray!

## Model Architecture

### Description

I feel the code below provides some of the relevant hyper-parameters used for tuning and the deep network architecture. This is basically a copy of the LeNet solution, with DropOut and regularization added being the main difference. In addition, one of the fully-connected layers was dropped. The other main difference is the depths for the convolution layers were upped a bit more as the LeNet solution was for 10 classes, and since we have 43, there are surely more finer-grained features to learn.

```
In [5]: EPOCHS = 500
        BATCH_SIZE = 1024 #orig 128
        # Arguments used for tf.truncated_normal, randomly defines variables f
        or the weights and biases for each layer
        mu = 0
        sigma = 0.16

        # Declaring weights outside method for access later. CLEANME: Move th
        is to a proper class
        depth_conv1 = 6
        depth_conv2 = 16
        size_fc_1 = 120
        size_fc_2 = 84

        conv1_W = tf.Variable(tf.truncated_normal(shape=(5, 5, 3, depth_conv1)
        , mean = mu, stddev = sigma))
        conv2_W = tf.Variable(tf.truncated_normal(shape=(5, 5, depth_conv1, de
        pth_conv2), mean = mu, stddev = sigma))
        fc1_W = tf.Variable(tf.truncated_normal(shape=(5*5*depth_conv2, size_f
        c_1), mean = mu, stddev = sigma))
        fc2_W = tf.Variable(tf.truncated_normal(shape=(size_fc_1, size_fc_2),
        mean = mu, stddev = sigma))
        fc3_W = tf.Variable(tf.truncated_normal(shape=(size_fc_2, num_classes
        ), mean = mu, stddev = sigma))
        keep_prob = tf.placeholder(tf.float32) # probability to keep units
```



```

def LeNet(x):
    print("Charging up!")

    # Layer 1: Convolutional. Input = 32x32x1. Output = 28x28xdepth_conv1.
    conv1_b = tf.Variable(tf.zeros(depth_conv1))
    conv1 = tf.nn.conv2d(x, conv1_W, strides=[1, 1, 1, 1], padding='VALID') + conv1_b

    # Activation.
    conv1 = tf.nn.relu(conv1)
    #conv1 = tf.nn.dropout(conv1, keep_prob)

    # Pooling. Input = 28x28xdepth_conv1. Output = 14x14xdepth_conv1.
    conv1 = tf.nn.max_pool(conv1, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='VALID')

    # Layer 2: Convolutional. Output = 10x10xdepth_conv2.
    conv2_b = tf.Variable(tf.zeros(depth_conv2))
    conv2 = tf.nn.conv2d(conv1, conv2_W, strides=[1, 1, 1, 1], padding='VALID') + conv2_b

    # Activation.
    conv2 = tf.nn.relu(conv2)
    #conv2 = tf.nn.dropout(conv2, keep_prob)

    # Pooling. Input = 10x10xdepth_conv2. Output = 5x5xdepth_conv2.
    conv2 = tf.nn.max_pool(conv2, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='VALID')

    # Flatten. Input = 5x5xdepth_conv2. Output = 1800.
    fc0 = flatten(conv2)

    # Layer 3: Fully Connected. Input = 1800. Output = size_fc_1.
    fc1_b = tf.Variable(tf.zeros(size_fc_1))
    fc1 = tf.matmul(fc0, fc1_W) + fc1_b

    # Activation.
    fc1 = tf.nn.relu(fc1)
    #fc1 = tf.nn.dropout(fc1, keep_prob)

    # Layer 3: Fully Connected. Input = size_fc_1. Output = size_fc_2.
    fc2_b = tf.Variable(tf.zeros(size_fc_2))
    fc2 = tf.matmul(fc1, fc2_W) + fc2_b

    # Activation.
    fc2 = tf.nn.relu(fc2)
    fc2 = tf.nn.dropout(fc2, keep_prob)

```

```

    # Layer 4: Fully Connected. Input = size_fc_2. Output = num_classes.
    fc3_b = tf.Variable(tf.zeros(num_classes))
    logits = tf.matmul(fc2, fc3_W) + fc3_b

    return logits

x = tf.placeholder(tf.float32, (None, 32, 32, 3))
y = tf.placeholder(tf.int32, (None))
one_hot_y = tf.one_hot(y, num_classes)

print('All done')

```

All done

## Train, Validate and Test the Model

```

In [21]: %%time

rate = 0.00175
largeBiasPenalty = 0.025
keep_probability = 0.79

logits = LeNet(x)
cross_entropy = tf.nn.softmax_cross_entropy_with_logits(labels=one_hot_y, logits=logits)

# Add regularization to the loss for highly-biased weights
loss_operation = tf.reduce_mean(cross_entropy) + \
    largeBiasPenalty * tf.nn.l2_loss(conv1_W) + \
    largeBiasPenalty * tf.nn.l2_loss(conv2_W) + \
    largeBiasPenalty * tf.nn.l2_loss(fc1_W) + \
    largeBiasPenalty * tf.nn.l2_loss(fc2_W) + \
    largeBiasPenalty * tf.nn.l2_loss(fc3_W)

optimizer = tf.train.AdamOptimizer(learning_rate = rate)
training_operation = optimizer.minimize(loss_operation)

### Calculate and report the accuracy on the training and validation sets.
correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(one_hot_y, 1))
accuracy_operation = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
saver = tf.train.Saver()

def evaluate(X_data, y_data):

```

```

    num_examples = len(X_data)
    total_accuracy = 0
    sess = tf.get_default_session()
    for offset in range(0, num_examples, BATCH_SIZE):
        batch_x, batch_y = X_data[offset:offset+BATCH_SIZE], y_data[offset:offset+BATCH_SIZE]
        accuracy = sess.run(accuracy_operation, feed_dict={x: batch_x, y: batch_y, keep_prob: 1.0})
        total_accuracy += (accuracy * len(batch_x))
    return total_accuracy / num_examples

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    num_examples = len(X_train)

    print("Training...")
    print()
    for i in range(EPOCHS):
        print("EPOCH {} ...".format(i+1))
        X_train, y_train = shuffle(X_train, y_train)
        for offset in range(0, num_examples, BATCH_SIZE):
            end = offset + BATCH_SIZE
            batch_x, batch_y = X_train[offset:end], y_train[offset:end]
        ]
        sess.run(training_operation, feed_dict={x: batch_x, y: batch_y, keep_prob: keep_probability})

        training_accuracy = evaluate(X_train, y_train)
        validation_accuracy = evaluate(X_validation, y_validation)
        print("{:.3f} = Training accuracy vs Validation Accuracy = {:.3f}".format(training_accuracy, validation_accuracy))

        print("Starting to save model...")
        saver.save(sess, './traffick')
        print("...model saved. Onward, hiya!")

    print()

print("ALL DONE! Yippie Kay Yae!")

```

Charging up!

Training...

EPOCH 1 ...

0.065 = Training accuracy vs Validation Accuracy = 0.061

Starting to save model...

...model saved. Onward, hiya!

EPOCH 2 ...  
0.055 = Training accuracy vs Validation Accuracy = 0.054  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 3 ...  
0.070 = Training accuracy vs Validation Accuracy = 0.068  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 4 ...  
0.080 = Training accuracy vs Validation Accuracy = 0.077  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 5 ...  
0.127 = Training accuracy vs Validation Accuracy = 0.120  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 6 ...  
0.144 = Training accuracy vs Validation Accuracy = 0.139  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 7 ...  
0.172 = Training accuracy vs Validation Accuracy = 0.169  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 8 ...  
0.218 = Training accuracy vs Validation Accuracy = 0.210  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 9 ...  
0.293 = Training accuracy vs Validation Accuracy = 0.283  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 10 ...  
0.380 = Training accuracy vs Validation Accuracy = 0.372  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 11 ...  
0.437 = Training accuracy vs Validation Accuracy = 0.425  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 12 ...  
0.480 = Training accuracy vs Validation Accuracy = 0.474  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 13 ...  
0.520 = Training accuracy vs Validation Accuracy = 0.511  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 14 ...  
0.532 = Training accuracy vs Validation Accuracy = 0.520  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 15 ...  
0.562 = Training accuracy vs Validation Accuracy = 0.549  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 16 ...  
0.562 = Training accuracy vs Validation Accuracy = 0.549  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 17 ...  
0.591 = Training accuracy vs Validation Accuracy = 0.578  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 18 ...  
0.614 = Training accuracy vs Validation Accuracy = 0.598  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 19 ...  
0.649 = Training accuracy vs Validation Accuracy = 0.637  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 20 ...  
0.666 = Training accuracy vs Validation Accuracy = 0.654  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 21 ...  
0.701 = Training accuracy vs Validation Accuracy = 0.693  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 22 ...  
0.701 = Training accuracy vs Validation Accuracy = 0.690  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 23 ...  
0.732 = Training accuracy vs Validation Accuracy = 0.718  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 24 ...  
0.737 = Training accuracy vs Validation Accuracy = 0.718  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 25 ...  
0.756 = Training accuracy vs Validation Accuracy = 0.744  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 26 ...  
0.781 = Training accuracy vs Validation Accuracy = 0.765  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 27 ...  
0.795 = Training accuracy vs Validation Accuracy = 0.783  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 28 ...  
0.795 = Training accuracy vs Validation Accuracy = 0.787  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 29 ...  
0.803 = Training accuracy vs Validation Accuracy = 0.790  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 30 ...  
0.805 = Training accuracy vs Validation Accuracy = 0.790  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 31 ...  
0.830 = Training accuracy vs Validation Accuracy = 0.816  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 32 ...  
0.823 = Training accuracy vs Validation Accuracy = 0.810  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 33 ...  
0.830 = Training accuracy vs Validation Accuracy = 0.811  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 34 ...  
0.843 = Training accuracy vs Validation Accuracy = 0.820  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 35 ...  
0.854 = Training accuracy vs Validation Accuracy = 0.835  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 36 ...  
0.843 = Training accuracy vs Validation Accuracy = 0.823  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 37 ...  
0.861 = Training accuracy vs Validation Accuracy = 0.840  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 38 ...  
0.872 = Training accuracy vs Validation Accuracy = 0.853  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 39 ...  
0.873 = Training accuracy vs Validation Accuracy = 0.856  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 40 ...  
0.874 = Training accuracy vs Validation Accuracy = 0.855  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 41 ...  
0.880 = Training accuracy vs Validation Accuracy = 0.859  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 42 ...  
0.884 = Training accuracy vs Validation Accuracy = 0.868  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 43 ...  
0.875 = Training accuracy vs Validation Accuracy = 0.854  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 44 ...  
0.897 = Training accuracy vs Validation Accuracy = 0.877  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 45 ...  
0.898 = Training accuracy vs Validation Accuracy = 0.879  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 46 ...  
0.910 = Training accuracy vs Validation Accuracy = 0.894  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 47 ...  
0.901 = Training accuracy vs Validation Accuracy = 0.888  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 48 ...  
0.910 = Training accuracy vs Validation Accuracy = 0.892  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 49 ...  
0.918 = Training accuracy vs Validation Accuracy = 0.903  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 50 ...  
0.918 = Training accuracy vs Validation Accuracy = 0.903  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 51 ...  
0.929 = Training accuracy vs Validation Accuracy = 0.914  
Starting to save model...  
...model saved. Onward, hiya!



EPOCH 52 ...  
0.924 = Training accuracy vs Validation Accuracy = 0.903  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 53 ...  
0.935 = Training accuracy vs Validation Accuracy = 0.915  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 54 ...  
0.935 = Training accuracy vs Validation Accuracy = 0.915  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 55 ...  
0.937 = Training accuracy vs Validation Accuracy = 0.918  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 56 ...  
0.941 = Training accuracy vs Validation Accuracy = 0.919  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 57 ...  
0.942 = Training accuracy vs Validation Accuracy = 0.921  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 58 ...  
0.943 = Training accuracy vs Validation Accuracy = 0.924  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 59 ...  
0.941 = Training accuracy vs Validation Accuracy = 0.923  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 60 ...  
0.946 = Training accuracy vs Validation Accuracy = 0.924  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 61 ...  
0.938 = Training accuracy vs Validation Accuracy = 0.922  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 62 ...  
0.953 = Training accuracy vs Validation Accuracy = 0.934  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 63 ...  
0.947 = Training accuracy vs Validation Accuracy = 0.929  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 64 ...  
0.949 = Training accuracy vs Validation Accuracy = 0.928  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 65 ...  
0.958 = Training accuracy vs Validation Accuracy = 0.936  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 66 ...  
0.956 = Training accuracy vs Validation Accuracy = 0.935  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 67 ...  
0.961 = Training accuracy vs Validation Accuracy = 0.938  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 68 ...  
0.954 = Training accuracy vs Validation Accuracy = 0.931  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 69 ...  
0.963 = Training accuracy vs Validation Accuracy = 0.940  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 70 ...  
0.959 = Training accuracy vs Validation Accuracy = 0.939  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 71 ...  
0.962 = Training accuracy vs Validation Accuracy = 0.938  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 72 ...  
0.962 = Training accuracy vs Validation Accuracy = 0.942  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 73 ...  
0.963 = Training accuracy vs Validation Accuracy = 0.943  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 74 ...  
0.966 = Training accuracy vs Validation Accuracy = 0.947  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 75 ...  
0.964 = Training accuracy vs Validation Accuracy = 0.941  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 76 ...  
0.955 = Training accuracy vs Validation Accuracy = 0.933  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 77 ...  
0.960 = Training accuracy vs Validation Accuracy = 0.942  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 78 ...  
0.972 = Training accuracy vs Validation Accuracy = 0.952  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 79 ...  
0.966 = Training accuracy vs Validation Accuracy = 0.946  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 80 ...  
0.968 = Training accuracy vs Validation Accuracy = 0.948  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 81 ...  
0.969 = Training accuracy vs Validation Accuracy = 0.947  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 82 ...  
0.970 = Training accuracy vs Validation Accuracy = 0.950  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 83 ...  
0.974 = Training accuracy vs Validation Accuracy = 0.955  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 84 ...  
0.972 = Training accuracy vs Validation Accuracy = 0.951  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 85 ...  
0.971 = Training accuracy vs Validation Accuracy = 0.949  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 86 ...  
0.971 = Training accuracy vs Validation Accuracy = 0.951  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 87 ...  
0.975 = Training accuracy vs Validation Accuracy = 0.954  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 88 ...  
0.966 = Training accuracy vs Validation Accuracy = 0.945  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 89 ...  
0.976 = Training accuracy vs Validation Accuracy = 0.959  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 90 ...  
0.976 = Training accuracy vs Validation Accuracy = 0.955  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 91 ...  
0.977 = Training accuracy vs Validation Accuracy = 0.959  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 92 ...  
0.968 = Training accuracy vs Validation Accuracy = 0.944  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 93 ...  
0.979 = Training accuracy vs Validation Accuracy = 0.959  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 94 ...  
0.975 = Training accuracy vs Validation Accuracy = 0.955  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 95 ...  
0.979 = Training accuracy vs Validation Accuracy = 0.960  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 96 ...  
0.975 = Training accuracy vs Validation Accuracy = 0.955  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 97 ...  
0.980 = Training accuracy vs Validation Accuracy = 0.959  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 98 ...  
0.977 = Training accuracy vs Validation Accuracy = 0.958  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 99 ...  
0.981 = Training accuracy vs Validation Accuracy = 0.963  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 100 ...  
0.977 = Training accuracy vs Validation Accuracy = 0.959  
Starting to save model...  
...model saved. Onward, hiya!

EPOCH 101 ...  
0.980 = Training accuracy vs Validation Accuracy = 0.961  
Starting to save model...  
...model saved. Onward, hiya!

```

EPOCH 102 ...
0.980 = Training accuracy vs Validation Accuracy = 0.963
Starting to save model...
...model saved.  Onward, hiya!

```

```

EPOCH 103 ...
0.980 = Training accuracy vs Validation Accuracy = 0.962
Starting to save model...
...model saved.  Onward, hiya!

```

```

EPOCH 104 ...

```

```

-----
-----
KeyboardInterrupt                                Traceback (most recent call
last)
<timed exec> in <module>()

<timed exec> in evaluate(X_data, y_data)

/home/carnd/anaconda3/envs/carnd-term1/lib/python3.5/site-packages/t
ensorflow/python/client/session.py in run(self, fetches, feed_dict,
options, run_metadata)
    764     try:
    765         result = self._run(None, fetches, feed_dict, options_p
tr,
--> 766                                     run_metadata_ptr)
    767     if run_metadata:
    768         proto_data = tf_session.TF_GetBuffer(run_metadata_ptr)

/home/carnd/anaconda3/envs/carnd-term1/lib/python3.5/site-packages/t
ensorflow/python/client/session.py in _run(self, handle, fetches, fe
ed_dict, options, run_metadata)
    962     if final_fetches or final_targets:
    963         results = self._do_run(handle, final_targets, final_fe
tches,
--> 964                                     feed_dict_string, options, run_
metadata)
    965     else:
    966         results = []

/home/carnd/anaconda3/envs/carnd-term1/lib/python3.5/site-packages/t
ensorflow/python/client/session.py in _do_run(self, handle, target_l
ist, fetch_list, feed_dict, options, run_metadata)
   1012     if handle is None:
   1013         return self._do_call(_run_fn, self._session, feed_dict
, fetch_list,
-> 1014                                     target_list, options, run_metadat
a)

```

```

1015         else:
1016             return self._do_call(_prun_fn, self._session, handle,
feed_dict,

/home/carnd/anaconda3/envs/carnd-term1/lib/python3.5/site-packages/t
ensorflow/python/client/session.py in _do_call(self, fn, *args)
1019     def _do_call(self, fn, *args):
1020         try:
-> 1021             return fn(*args)
1022         except errors.OpError as e:
1023             message = compat.as_text(e.message)

/home/carnd/anaconda3/envs/carnd-term1/lib/python3.5/site-packages/t
ensorflow/python/client/session.py in _run_fn(session, feed_dict, fe
tch_list, target_list, options, run_metadata)
1001         return tf_session.TF_Run(session, options,
1002                                     feed_dict, fetch_list,
target_list,
-> 1003                                     status, run_metadata)
1004
1005     def _prun_fn(session, handle, feed_dict, fetch_list):

KeyboardInterrupt:

```

And what do we get!!?? PASS!!!! Kinda...

```

In [22]: with tf.Session() as sess:
saver.restore(sess, tf.train.latest_checkpoint('.'))

test_accuracy = evaluate(X_test, y_test)
print("Test Accuracy = {:.3f}".format(test_accuracy))

```

Test Accuracy = 0.882

Below preserves the one time I actually got the required accuracy. Yes, its dependent on how the dataset is shuffled, but I have spent too long on this project, need to catch up with the rest of the class!

```
In [133]: %%time

with tf.Session() as sess:
    saver.restore(sess, tf.train.latest_checkpoint('.'))

    test_accuracy = evaluate(X_test, y_test)
    print("Test Accuracy = {:.3f}".format(test_accuracy))
```

```
Test Accuracy = 0.932
CPU times: user 16.3 s, sys: 3.6 s, total: 19.9 s
Wall time: 15.4 s
```

## Step 4: Use the model to make predictions on new images

Here we download 5 images from the web to test our dataset against. I added a sixth image. For some silly reason when I first started, I sought out traffic signs that are *not* in the dataset (we, afterall, have a test set for verifying overall accuracy anyway). So I included a sixth one, a Tank-cross sign, which doesn't exist in the dataset, to see how easily the model may be fooled.

### Load and Output the Images

```
In [27]: # class id = NONE!
tank = mpimg.imread('downloaded/tank.png')
print('size of tank: {}'.format(np.shape(tank)))

# It seems like mpimg is already 'normalizing' the images as we read t
hem. Or it was OSX Preview as I reshaped the classes.
#print('full output of tank: {}'.format(tank))
#print('normalizing tank: {}'.format(normalize(tank)))

plt.title('TANK!')
plt.imshow(tank)
plt.show()

# class id = 31
wild_animals = mpimg.imread('downloaded/847px-Zeichen_142-10_-_Wildwec
hsel,_Aufstellung_rechts,_StVO_1992.svg.png')
print('size of animal: {}'.format(np.shape(wild_animals)))
plt.title('Wild Animal Xing')
plt.imshow(wild_animals)
plt.show()
```



```

# class id = 18
general_caution = mpimg.imread('downloaded/Zeichen_101_-_Gefahrstelle,
_StVO_1970.svg.png')
print('size of general_caution: {}'.format(np.shape(general_caution))
)
plt.title('General Caution')
plt.imshow(general_caution)
plt.show()

# class id = 19
dangerous_left = mpimg.imread('downloaded/Zeichen_103-10_-_Kurve_(link
s),_StVO_1992.svg.png')
print('size of dangerous_left: {}'.format(np.shape(dangerous_left)))
plt.title('Dangerous curve to the left')
plt.imshow(dangerous_left)
plt.show()

# class id = 3
kph_60 = mpimg.imread('downloaded/Zeichen_274-56.svg.png')
print('size of kph_60: {}'.format(np.shape(kph_60)))
plt.title('60 KM/h')
plt.imshow(kph_60)
plt.show()

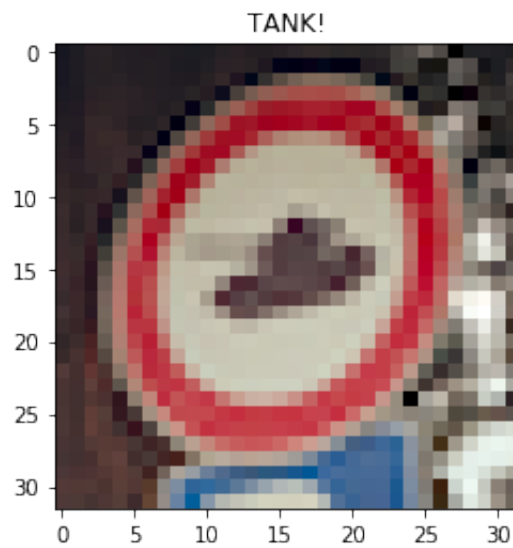
# class id 1
kph_30 = mpimg.imread('downloaded/Zeichen_274.1_-_Beginn_einer_Tempo_3
0-Zone,_StVO_2013.svg.png')
print('size of kph_30: {}'.format(np.shape(kph_30)))
plt.title('Entering 30 KM/h zone')
plt.imshow(kph_30)
plt.show()

X_downloaded = [tank, wild_animals, general_caution, dangerous_left, k
ph_60, kph_30]
y_downloaded = [-1,31,18,19,3,1]

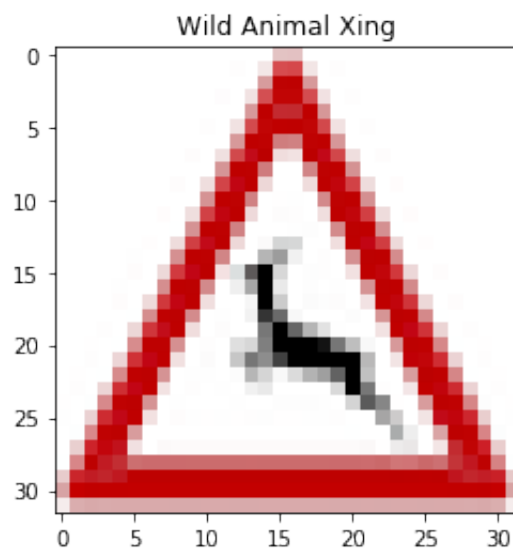
print(np.shape(X_downloaded))

size of tank: (32, 32, 3)

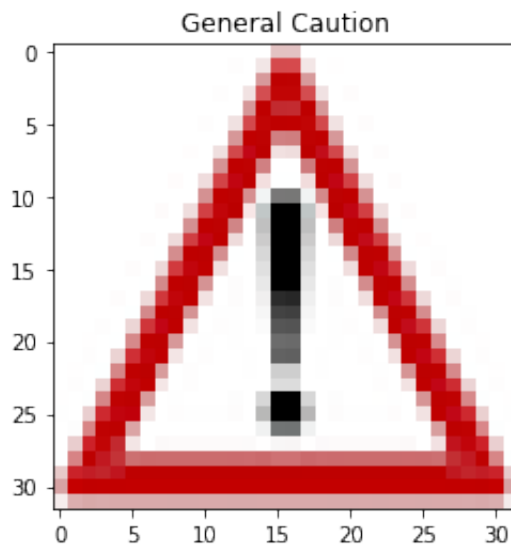
```



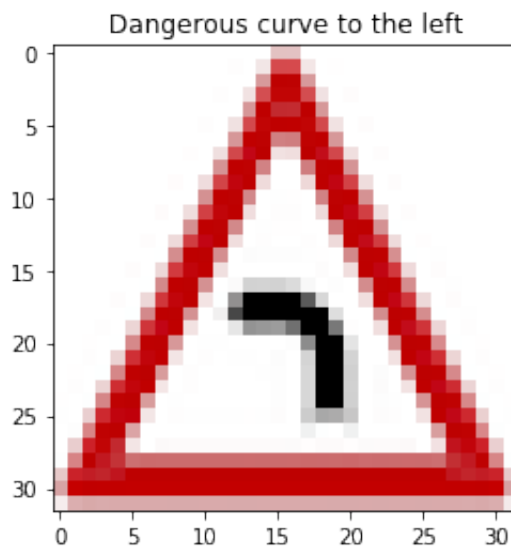
size of animal: (32, 32, 3)



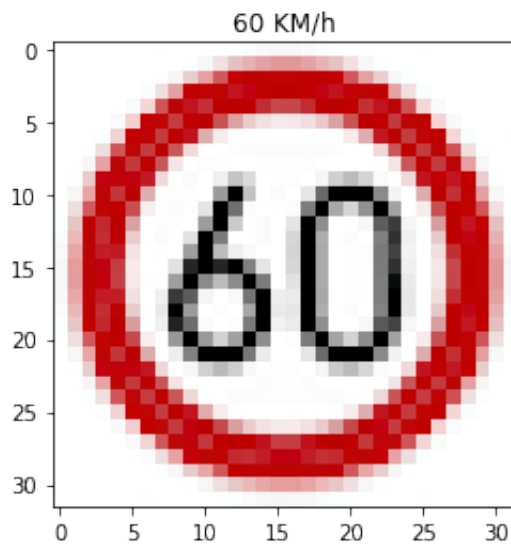
size of general\_caution: (32, 32, 3)



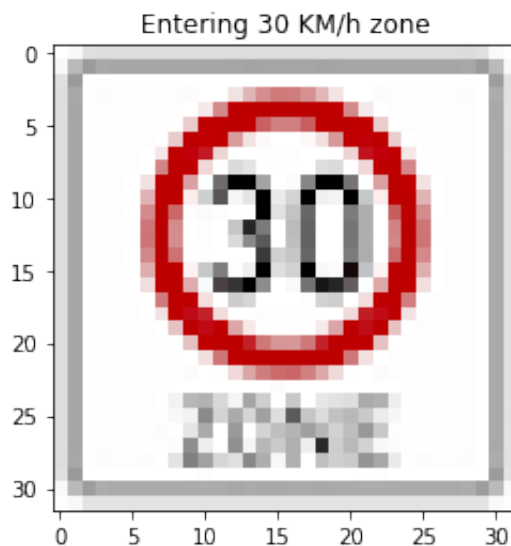
size of dangerous\_left: (32, 32, 3)



size of kph\_60: (32, 32, 3)



size of kph\_30: (32, 32, 3)



(6, 32, 32, 3)

---

## Step 5 Analyze the softmax probabilities of the new images Test a Model on New Images

### Predict the Sign Type for Each Image

In [23]: %%time

```
with tf.Session() as sess:
    saver.restore(sess, tf.train.latest_checkpoint('.'))
    test_accuracy = evaluate(X_downloaded, y_downloaded)
    print("Test Accuracy = {:.3f}".format(test_accuracy))
```

Test Accuracy = 0.000

CPU times: user 564 ms, sys: 460 ms, total: 1.02 s

Wall time: 528 ms

## Output Top 5 Softmax Probabilities For Each Image Found on the Web

In [24]: *# function to print out softmax*  
*# It may be a bit more optimized to run this during accuracy option. But not needed for this project's purposes*

```
top5 = tf.nn.top_k(tf.nn.softmax(logits), k=5)
```

```
def eval_downloaded(X_data, y_data):
    num_examples = len(X_data)
    sess = tf.get_default_session()
    t5 = sess.run(top5, feed_dict={x: X_data, y: y_data, keep_prob: 1.0})
    print("TOP5    {}".format(t5))
    return t5
```

```
print('Eval with top5 function loaded')
```

Eval with top5 function loaded

In [25]: %%time

```
with tf.Session() as sess:
    saver.restore(sess, tf.train.latest_checkpoint('.'))
    top5_accuracy = eval_downloaded(X_downloaded, y_downloaded)
    print("top 5 matrix in full = {}".format(top5_accuracy))
```

```

TOP5    TopKV2(values=array([[ 0.0981414 ,  0.07578082,  0.06946859,
0.06201295,  0.04533249],
      [ 0.09399547,  0.08496545,  0.06298173,  0.06221561,  0.04303
629],
      [ 0.08181851,  0.07178923,  0.07077026,  0.06547637,  0.04611
111],
      [ 0.10083228,  0.08288508,  0.07090969,  0.05555868,  0.04724
773],
      [ 0.09955632,  0.08115049,  0.08067164,  0.06269072,  0.04845
662],
      [ 0.08280152,  0.07337368,  0.07168196,  0.0580777 ,  0.04537
925]]), dtype=float32), indices=array([[10,  5,  3,  8,  9],
      [10,  5,  8,  3, 25],
      [ 5,  8, 10,  3, 20],
      [10,  5,  3,  8, 20],
      [ 5,  3, 10,  8,  4],
      [10,  5,  3,  8, 20]], dtype=int32))
top 5 matrix in full = TopKV2(values=array([[ 0.0981414 ,  0.0757808
2,  0.06946859,  0.06201295,  0.04533249],
      [ 0.09399547,  0.08496545,  0.06298173,  0.06221561,  0.04303
629],
      [ 0.08181851,  0.07178923,  0.07077026,  0.06547637,  0.04611
111],
      [ 0.10083228,  0.08288508,  0.07090969,  0.05555868,  0.04724
773],
      [ 0.09955632,  0.08115049,  0.08067164,  0.06269072,  0.04845
662],
      [ 0.08280152,  0.07337368,  0.07168196,  0.0580777 ,  0.04537
925]]), dtype=float32), indices=array([[10,  5,  3,  8,  9],
      [10,  5,  8,  3, 25],
      [ 5,  8, 10,  3, 20],
      [10,  5,  3,  8, 20],
      [ 5,  3, 10,  8,  4],
      [10,  5,  3,  8, 20]], dtype=int32))
CPU times: user 564 ms, sys: 516 ms, total: 1.08 s
Wall time: 545 ms

```

Overall, these numbers are not surprising. Whatever preprocessing steps I did on the images clearly did not go over so well. Given more time (must move on to next project), I am confident I can get this to be more interesting. At least I can show that I know enough Tensorflow to evaluate top-5. What would be interesting would be one would expect to hope to see low probabilities for a completely unknown image. This would maybe also require more augmentation.

## Step 6 Project Writeup

This here Jupyter notebook acts as my writeup!

Overall, I had a frustrating time on this project, as fun as it was to play around with tensorflow and learn actual traffic signs. Due to work obligations, I could realistically only commit an hour or two per day to this project. The majority of my time I spent spinning my wheels due to trying to get the saved model to restore from the saved checkpoint. Using 'last checkpoint' works fine, but the other approaches as documented in the TensorFlow documentation and Udacity class only resulted in errors.

I also was surprised that the LeNet architecture worked perfectly fine, even without normalization and in spite of the fact it had been setup for 10-class problem, rather than the 43-class problem we currently have. Also surprising is adding dropout dramatically decreased the accuracy. I noticed that, in spite of a lot of parameters, there seemed to be a local optima that kept accuracy stuck at exactly 5.4%. Sometimes, I would let the model run for 70 epochs and then it would rapidly get close to 95% accuracy after getting out of the 5.4% rut it was in.

I do not agree with augmenting the dataset, despite explicit hints provided, and feel it's important to construct a deep net that realistically simulates learning conditions of the real world. Even better would be to use few examples. The average human, in their lifetime, would not need to see 100,000+ examples of traffic signs repeatedly in order to accurately predict them. Granted, we are also in these training sets teaching a net to read images. It would be nice if there were preconfigured nets that already understand how to read images, that then feed into one specifically for traffic signs. Some of my comments are inspired by a tweet from Yan Lecunn, I'll let it speak for itself when he was congratulating Alpha Go on its achievements, and then challenged them to do the same with far less examples and reinforcement learning. [Yann's tweet \(https://twitter.com/ylecun/status/708732919548919808\)](https://twitter.com/ylecun/status/708732919548919808)

In either case I was able to get the minimum necessary accuracy.

Some other observations:

- I tried to apply the "keep calm and lower your learning rate" trick. I was giving values as low as a magnitude of  $10e-10$  and I still constantly kept getting stuck in the 5% rut. I reread what the Adam optimizer is and realize it applies the learn-rate lowering throughout repeated iterations, so I embraced a higher learning rate and that did a much better job of not being stuck in the 5% gutter.
- I removed a fully-connected layer. It took a lot more cajoling to get results above 90%, so I brought it back.
- I used a lazy-person's stop for overregularization. By saving the model after each epoch, I could keep my eye on when it reached a satisfactory accuracy, and stop as needed. I thought about applying a simple-moving average to only stop when the rate of accuracy decreases after X epochs, but since I wasn't sure what values to use, and sometimes it would take dozens of minutes to reach a certain state, I would hate for things to stop. So I used an overly large number of epochs, with manual stopping.
- Strangely, dropout and regularization seemed to make things worst